

## Stage 4 - Entity Matching

### 4.1 Readin data and previous save matcher

```
In [1]: from collections import Counter
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import pandas as pd
import py_entitymatching as em
```

```
In [2]: data_dir = './dataset/structured_data/'
A_filename = data_dir+'forbes_all_rename.csv'
B_filename = data_dir+'nasdaq_rename.csv'
blocked_res = data_dir+'blocking_results.csv'
shared_fields = ['Name','Country', 'Industry', "MarketValue"] ## These
are shared fileds
# use the random forest matcher
m = em.load_object(data_dir+"rf_matcher.pkl"); # load matcher from pre
vious matching stage.
```

```
In [3]: # Load orignial tables
A = pd.read_csv(A_filename, encoding = "ISO-8859-1" );
em.set_key(A, 'id');
A1 = A[['id'] + shared_fields];
em.set_key(A1, 'id');
B = pd.read_csv(B_filename, encoding = "ISO-8859-1");
em.set_key(B, 'id');
B1 = B[['id'] + shared_fields];
em.set_key(B1, 'id');
# Load the pre-labeled data
S1 = em.read_csv_metadata(blocked_res,
                           key='_id',
                           ltable=A1, rtable=B1,
                           fk_ltable='ltable_id', fk_rtable='rtable_id',
                           encoding = "ISO-8859-1");
```

Metadata file is not present in the given path; proceeding to read the csv file.

## 4.2 Entity matching

Build features.

```
In [4]: # Generate a set of features
F = em.get_features_for_matching(A1, B1)
# Add some new feature to F
def MarketValue_ratio(ltuple, rtuple) :
    try :
        return float(ltuple.MarketValue) / float(rtuple.MarketValue)
    except ValueError :
        return 0
em.add_blackbox_feature(F, 'MarketValue_ratio', MarketValue_ratio)
```

Out[4]: True

```
In [5]: F.feature_name
```

```

Out[5]: 0          id_id_exm
        1          id_id_anm
        2          id_id_lev_dist
        3          id_id_lev_sim
        4      Name_Name_jac_qgm_3_qgm_3
        5      Name_Name_cos_dlm_dc0_dlm_dc0
        6      Name_Name_jac_dlm_dc0_dlm_dc0
        7          Name_Name_mel
        8      Name_Name_lev_dist
        9      Name_Name_lev_sim
       10      Name_Name_nmw
       11      Name_Name_sw
       12      Country_Country_jac_qgm_3_qgm_3
       13      Country_Country_cos_dlm_dc0_dlm_dc0
       14      Country_Country_jac_dlm_dc0_dlm_dc0
       15          Country_Country_mel
       16      Country_Country_lev_dist
       17      Country_Country_lev_sim
       18      Country_Country_nmw
       19      Country_Country_sw
       20      Industry_Industry_jac_qgm_3_qgm_3
       21      Industry_Industry_cos_dlm_dc0_dlm_dc0
       22      Industry_Industry_jac_dlm_dc0_dlm_dc0
       23          Industry_Industry_mel
       24      Industry_Industry_lev_dist
       25      Industry_Industry_lev_sim
       26      Industry_Industry_nmw
       27      Industry_Industry_sw
       28          MarketValue_ratio
Name: feature_name, dtype: object

```

Compute features for all blocked tuple pairs.

```

In [6]: L1 = em.extract_feature_vecs(S1, feature_table=F, attrs_after='is_matc
        h', show_progress=False)
        L1 = em.impute_table(L1,
                             exclude_attrs=['_id', 'ltable_id', 'rtable_id', "is
        _match"],
                             strategy='mean')

```

Predict using the matcher.

```

In [7]: predictions = m.predict(table=L1, exclude_attrs=['_id', 'ltable_id', '
        rtable_id', "is_match"],
        append=True, target_attr='predicted', inplace=False)

```

```
In [8]: predictions.to_csv(data_dir+"matching_result.csv") # save matching results to data directory
```

## 4.3 Find duplicates

```
In [9]: import collections
fd = collections.defaultdict(list) # dictionary for forbes
nd = collections.defaultdict(list) # dictionary for nasdaq
lid = predictions["ltable_id"]
rid = predictions["rtable_id"]
match = predictions["predicted"]

# forbes dictionary
for i in range(len(lid)):
    if match[i] == 0:
        continue
    fd[lid[i]].append(rid[i])
    nd[rid[i]].append(lid[i])
```

For the multiple matches for one entity, we checked the results manually and figured such similar entities are difficult to identify even by human beings. We printed one example below.

```
In [10]: S1[198:256]
```

```
Out[10]:
```

	Unnamed: 0	_id	ltable_id	rtable_id	ltable_Country	ltable_Industry	ltable_Market
198	4327	4327	1907	556	United States	Investment Services	60400
199	4344	4344	1907	557	United States	Investment Services	60400
200	4355	4355	1907	558	United States	Investment Services	60400
						Investment	

<b>201</b>	4366	4366	1907	559	United States	Services	60400
<b>202</b>	4371	4371	1907	560	United States	Investment Services	60400
<b>203</b>	4381	4381	1907	561	United States	Investment Services	60400
<b>204</b>	4415	4415	1907	562	United States	Investment Services	60400
<b>205</b>	4431	4431	1907	563	United States	Investment Services	60400
<b>206</b>	4438	4438	1907	564	United States	Investment Services	60400
<b>207</b>	4439	4439	1907	565	United States	Investment Services	60400
<b>208</b>	4440	4440	1907	566	United States	Investment Services	60400
<b>209</b>	4456	4456	1907	567	United States	Investment Services	60400
<b>210</b>	4471	4471	1907	568	United States	Investment Services	60400
						Investment	

<b>211</b>	4503	4503	1907	569	United States	Services	60400
<b>212</b>	4520	4520	1907	570	United States	Investment Services	60400
<b>213</b>	4537	4537	1907	571	United States	Investment Services	60400
<b>214</b>	4633	4633	1907	573	United States	Investment Services	60400
<b>215</b>	4648	4648	1907	574	United States	Investment Services	60400
<b>216</b>	4660	4660	1907	575	United States	Investment Services	60400
<b>217</b>	4678	4678	1907	576	United States	Investment Services	60400
<b>218</b>	4682	4682	1907	577	United States	Investment Services	60400
<b>219</b>	4692	4692	1907	578	United States	Investment Services	60400
						Investment	

<b>220</b>	4698	4698	1907	579	United States	Services	60400
<b>221</b>	4699	4699	1907	580	United States	Investment Services	60400
<b>222</b>	4710	4710	1907	581	United States	Investment Services	60400
<b>223</b>	4722	4722	1907	582	United States	Investment Services	60400
<b>224</b>	4742	4742	1907	583	United States	Investment Services	60400
<b>225</b>	4759	4759	1907	584	United States	Investment Services	60400
<b>226</b>	4770	4770	1907	585	United States	Investment Services	60400
<b>227</b>	4781	4781	1907	586	United States	Investment Services	60400
<b>228</b>	4799	4799	1907	587	United States	Investment Services	60400
<b>229</b>	4803	4803	1907	588	United States	Investment Services	60400
						Investment	

<b>230</b>	4804	4804	1907	589	United States	Services	60400
<b>231</b>	4805	4805	1907	590	United States	Investment Services	60400
<b>232</b>	4806	4806	1907	591	United States	Investment Services	60400
<b>233</b>	4820	4820	1907	592	United States	Investment Services	60400
<b>234</b>	4833	4833	1907	593	United States	Investment Services	60400
<b>235</b>	4848	4848	1907	594	United States	Investment Services	60400
<b>236</b>	4854	4854	1907	595	United States	Investment Services	60400
<b>237</b>	4855	4855	1907	596	United States	Investment Services	60400
<b>238</b>	4856	4856	1907	597	United States	Investment Services	60400
<b>239</b>	4857	4857	1907	598	United States	Investment Services	60400



<b>240</b>	4858	4858	1907	599	United States	Investment Services	60400
<b>241</b>	4859	4859	1907	600	United States	Investment Services	60400
<b>242</b>	4873	4873	1907	601	United States	Investment Services	60400
<b>243</b>	4889	4889	1907	602	United States	Investment Services	60400
<b>244</b>	4892	4892	1907	603	United States	Investment Services	60400
<b>245</b>	4903	4903	1907	604	United States	Investment Services	60400
<b>246</b>	4918	4918	1907	605	United States	Investment Services	60400
<b>247</b>	4924	4924	1907	606	United States	Investment Services	60400
<b>248</b>	4925	4925	1907	607	United States	Investment Services	60400

<b>249</b>	4926	4926	1907	608	United States	Investment Services	60400
<b>250</b>	4927	4927	1907	609	United States	Investment Services	60400
<b>251</b>	4949	4949	1907	610	United States	Investment Services	60400
<b>252</b>	4975	4975	1907	611	United States	Investment Services	60400
<b>253</b>	4997	4997	1907	612	United States	Investment Services	60400
<b>254</b>	5040	5040	1907	613	United States	Investment Services	60400
<b>255</b>	5065	5065	1907	614	United States	Investment Services	60400

Thus, we choose to only care about one-to-one matches, resulting in a total number of 704 matched entities.

```
In [11]: ## list of one-to-one matches
total_match = list() # list(l_id, r_id))

for ltable_id in fd :
    if len(fd[ltable_id]) == 1 :
        rtable_id = fd[ltable_id][0]
        if len(nd[rtable_id]) == 1 :
            total_match.append((ltable_id, rtable_id))

print (len(total_match))
```

703

## 4.4 New schema for combined table

First, we check the existing schema.

- Fields shared with both tables

```
In [12]: shared_fields
```

```
Out[12]: ['Name', 'Country', 'Industry', 'MarketValue']
```

- Fields only in table A (Forbes)

```
In [13]: A_unique_fields = list(set(A.columns) - set(B.columns))
A_unique_fields
```

```
Out[13]: ['Assets', 'Employee', 'Sales', 'Profits']
```

- Fields only in table B (NASDAQ)

```
In [14]: B_unique_fields = list(set(B.columns) - set(A.columns))
B_unique_fields
```

```
Out[14]: ['IPOyear', 'Symbol', 'LastSale', 'Summary Quote', 'Sector']
```

Based on the features in two tables, we define the new table schema to be the combine of all fields above, plus 'ltable\_id' and 'rtable\_id' as foreign key to original tables A and B.

```
In [15]: combined_fields = shared_fields + A_unique_fields + B_unique_fields
E_fields = ['ltable_id', 'rtable_id'] + combined_fields
E_fields
```

```
Out[15]: ['ltable_id',
          'rtable_id',
          'Name',
          'Country',
          'Industry',
          'MarketValue',
          'Assets',
          'Employee',
          'Sales',
          'Profits',
          'IPOyear',
          'Symbol',
          'LastSale',
          'Summary Quote',
          'Sector']
```

## 4.5 Combining

Because NASDAQ table contains the company information formally registered on file, we use feature values from NASDAQ table if there is a conflict between two tables over shared fields (company name, industry, country).

```
In [16]: A_fields = shared_fields + A_unique_fields
B_fields = shared_fields + B_unique_fields
```

```
In [17]: A_indexed = A.set_index('id')
B_indexed = B.set_index('id')
```

```
In [36]: E_match = pd.DataFrame(index=range(len(total_match)), columns=E_fields
)
for i, (ltable_id, rtable_id) in enumerate(total_match) :
    E_match.loc[i, 'ltable_id'] = ltable_id
    E_match.loc[i, 'rtable_id'] = rtable_id
    E_match.loc[i] = E_match.loc[i].combine_first(B_indexed.ix[rtable_id][B_fields]).combine_first(A_indexed.ix[ltable_id][A_fields])

E_match.head()
```

Out[36]:

	ltable_id	rtable_id	Name	Country	Industry	MarketValue	Ass
0	2051	3485	Principal Financial Group Inc	United States	Accident &Health Insurance	18142.7	218
1	2053	3517	Prudential Financial, Inc.	United States	Life Insurance	45911.9	757
2	684	1103	Constellation Brands Inc	United States	Beverages (Production/Distribution)	31775.4	170
3	2061	4234	Torchmark Corporation	United States	Life Insurance	9082.82	199
4	19	372	Associated Banc-Corp	United States	Major Banks	3718.84	282

```

In [37]: unique_ltable_ids = [i for i in A.id if i not in set(ltable_id for ltable_id, _ in total_match)]
E_A_only = pd.DataFrame(index=range(len(unique_ltable_ids)), columns=E_fields)
for i, ltable_id in enumerate(unique_ltable_ids) :
    E_A_only.loc[i, 'ltable_id'] = ltable_id
    E_A_only.loc[i, A_fields] = A_indexed.ix[ltable_id][A_fields]

E_A_only.head()

```

Out[37]:

	ltable_id	rtable_id	Name	Country	Industry	MarketValue	Assets	Employee
0	1	NaN	77 Bank	Japan	Banks	1400	69100	NaN
1	2	NaN	Abu Dhabi Commercial Bank	United Arab Emirates	Banks	11000	62100	NaN
2	3	NaN	Abu Dhabi Islamic Bank	United Arab Emirates	Banks	3800	24300	NaN
3	4	NaN	Agricultural Bank of China	China	Banks	152700	2739800	NaN
4	5	NaN	Ahli United Bank	Bahrain	Banks	4200	34000	NaN

```
In [38]: unique_rtable_ids = [i for i in B.id if i not in set(rtable_id for _,
rtable_id in total_match)]
E_B_only = pd.DataFrame(index=range(len(unique_rtable_ids)), columns=E
_fields)
for i, rtable_id in enumerate(unique_rtable_ids) :
    E_B_only.loc[i, 'rtable_id'] = rtable_id
    E_B_only.loc[i, B_fields] = B_indexed.ix[rtable_id][B_fields]

E_B_only.head()
```

Out[38]:

	ltable_id	rtable_id	Name	Country	Industry	MarketValue	Assets
0	NaN	1	1-800 FLOWERS.COM, Inc.	United States	Other Specialty Stores	665.526	NaN
1	NaN	2	1347 Property Insurance Holdings, Inc.	United States	Property- Casualty Insurers	48.8455	NaN
2	NaN	3	180 Degree Capital Corp.	United States	Finance/Investors Services	44.8111	NaN
3	NaN	4	1st Constitution Bancorp (NJ)	United States	Savings Institutions	148.506	NaN
4	NaN	5	1st Source Corporation	United States	Major Banks	1262.61	NaN

```
In [39]: E = E_match.append(E_A_only).append(E_B_only).reset_index(drop=True)
         E.tail()
```

```
Out[39]:
```

	ltable_id	rtable_id	Name	Country	Industry	MarketVa
<b>7116</b>	NaN	4710	ZTO Express (Cayman) Inc.	China	Trucking Freight/Courier Services	9574.11
<b>7117</b>	NaN	4711	Zumiez Inc.	United States	Clothing/Shoe/Accessory Stores	456.463
<b>7118</b>	NaN	4712	Zweig Fund, Inc. (The)	United States	NaN	180.265
<b>7119</b>	NaN	4713	Zynerba Pharmaceuticals, Inc.	United States	Major Pharmaceuticals	265.618
<b>7120</b>	NaN	4714	Zynga Inc.	United States	EDP Services	2474.38

```
In [40]: E.to_csv(data_dir + 'E.csv')
```

## 4.6 Statistics of Table E

- Schema

```
In [41]: E.columns
```

```
Out[41]: Index(['ltable_id', 'rtable_id', 'Name', 'Country', 'Industry', 'MarketValue',
               'Assets', 'Employee', 'Sales', 'Profits', 'IPOyear', 'Symbol',
               'LastSale', 'Summary Quote', 'Sector'],
              dtype='object')
```

- Number of tuples

```
In [42]: len(E)
```

```
Out[42]: 7121
```



- Example tuples

1. Tuples from both tables A and B
2. Tuples from only table A
3. Tuples from only Table B

In [43]: E.loc[[1, 563]]

Out[43]:

	ltable_id	rtable_id	Name	Country	Industry	MarketValue	Assets	Employee
<b>1</b>	2053	3517	Prudential Financial, Inc.	United States	Life Insurance	45911.9	757400	NaN
<b>563</b>	1508	1930	Golub Capital BDC, Inc.	United States	Food Markets	1097.56	NaN	20626

In [44]: E.loc[[1000, 1234]]

Out[44]:

	ltable_id	rtable_id	Name	Country	Industry	MarketValue	Assets	Er
<b>1000</b>	316	NaN	China Biologic Products	Hong Kong	NaN	2941	NaN	NaN
<b>1234</b>	588	NaN	Inditex	Spain	Apparel/Footwear Retail	103200	18800	NaN

In [46]: E.loc[[5000, 7000]]

Out[46]:

	ltable_id	rtable_id	Name	Country	Industry	MarketValue	Assets	Er
<b>5000</b>	NaN	2217	Innoviva, Inc.	United States	Major Pharmaceuticals	1509.16	NaN	NaN
<b>7000</b>	NaN	4570	Western Asset Managed Municipals Fund, Inc.	United States	NaN	592.45	NaN	NaN

## 4.7 Merging function explained

We use NASDAQ as major table and Forbes as minor table. That is, we will priorly use values from table B for shared fields if the field is present in both tables.

This can be best seen with those matches whose shared fields have N/A in both of the original tables A and B. We list match ids here.

```
In [23]: for i_match, (ltable_id, rtable_id) in enumerate(total_match) :
          ltuple, rtuple = A_indexed.ix[ltable_id], B_indexed.ix[rtable_id]
          if any(ltuple[shared_fields].isnull()) and any(rtuple[shared_fields].isnull()) :
              print(i_match)
```

563

Here is a pair of tuples from A and B.

```
In [24]: i_match = 563
          total_match[i_match]
```

Out[24]: (1508, 1930)

```
In [25]: ltable_id, rtable_id = total_match[i_match]
          ltuple, rtuple = A_indexed.ix[ltable_id], B_indexed.ix[rtable_id]
```

```
In [26]: ltuple.ix[A_fields]
```

```
Out[26]: Name                Golub
          Country            United States
          Industry            Food Markets
          MarketValue         NaN
          Assets              NaN
          Employee            20626
          Sales               3400
          Profits             NaN
          Name: 1508, dtype: object
```

```
In [27]: rtuple.ix[B_fields]
```

```
Out[27]: Name                Golub Capital BDC, Inc.
        Country                United States
        Industry                NaN
        MarketValue            1097.56
        IPOyear                2010
        Symbol                GBDC
        LastSale                19.87
        Summary Quote    http://www.nasdaq.com/symbol/gbdc
        Sector                NaN
        Name: 1930, dtype: object
```

```
In [28]: combined_tuple = pd.Series(index=E_fields)
```

```
In [29]: combined_tuple = combined_tuple.combine_first(rtuple[B_fields])
        combined_tuple
```

```
Out[29]: Assets                NaN
        Country                United States
        Employee                NaN
        IPOyear                2010
        Industry                NaN
        LastSale                19.87
        MarketValue            1097.56
        Name                Golub Capital BDC, Inc.
        Profits                NaN
        Sales                NaN
        Sector                NaN
        Summary Quote    http://www.nasdaq.com/symbol/gbdc
        Symbol                GBDC
        ltable_id            NaN
        rtable_id            NaN
        dtype: object
```

```
In [30]: combined_tuple = combined_tuple.combine_first(ltuple[A_fields])
combined_tuple
```

```
Out[30]: Assets                NaN
Country                United States
Employee                20626
IPOyear                2010
Industry                Food Markets
LastSale                19.87
MarketValue            1097.56
Name                    Golub Capital BDC, Inc.
Profits                NaN
Sales                  3400
Sector                NaN
Summary Quote          http://www.nasdaq.com/symbol/gbdc
Symbol                GBDC
ltable_id              NaN
rtable_id              NaN
dtype: object
```

Note that we use value from table A for "Industry" which only presents in A, value from table B for "MarketValue" which only presents in B, and value from table B for "Name" which presents in both A and B.

This is exactly what we got in E.

```
In [31]: E_match.loc[i_match]
```

```
Out[31]: ltable_id              1508
rtable_id              1930
Name                    Golub Capital BDC, Inc.
Country                United States
Industry                Food Markets
MarketValue            1097.56
Assets                NaN
Employee                20626
Sales                  3400
Profits                NaN
IPOyear                2010
Symbol                GBDC
LastSale                19.87
Summary Quote          http://www.nasdaq.com/symbol/gbdc
Sector                NaN
Name: 563, dtype: object
```